

CS585 Final

Spring 2018: 5/3/18

Duration: 1 hour

Instructions/notes

- the exam is closed books/notes/devices/neighbors, and open mind :)
- there are 10 questions, plus a bonus
- there are no 'trick' questions
- please do NOT cheat; you get a 0 if you are found to have cheated
- when time is up, stop your work; you get a 0 if you continue
- good luck, hope you do well!

Q	Your score	Max possible score
1		2
2		4
3		4
4		4
5		2
6		3
7		2
8		4
9		5
10		5
Bonus		1
Total		36

Q1 (1+1=2 points). The usual 'MapReduce' steps are mapping, shuffling and reducing. But sometimes, an extra 'combining' step is inserted.

a. where does this occur (in the sequence of steps)?

b. what is its purpose (why include it)?

Q2 (2+2=4 points). You are running low on battery, driving around in an unfamiliar neighborhood - you need to recharge soon. You ask your navigation software for the nearest electric charging station, it instantly obliges - an R-tree helped, behind the scenes. Draw a simple map that shows your location, a scattering of 8 charging stations in the surrounding area, and an R-tree to contain the 8+1=9 locations; explain how the R-tree helps in quick retrieval of the result.

Q3 (4 points). In standard 'basket analysis', we search for rules of the type $A \rightarrow B$, where we specify 'support' threshold for A and a confidence for the \rightarrow . In other words, what are the (A,B) pairs such that given that A occurs, B occurs as well.

In 'differential basket analysis', we can take this a big step further - we can compare the occurrence of (A,B) pairs, given other (unrelated to (A,B)) factors. This can help make better use of the (A,B) associations. Eg. does (A,B) only occur (or NOT occur) in a certain store or groups of them?

What other factors (than store locations) can you think of? You can assume (bet!) that vast amount of customer data is available. List two factors related to customers, and two factors not related to customers.

1.

2.

3.

4.

Q4 (4 points). Before NoSQL, almost all databases were based on the relational model: tables (entities), and PK/FK relationships between them. The NoSQL paradigm offers us alternatives. There is one specific 'freedom' that NoSQL offers, from an architectural standpoint (ie. in the design, and redesign, of a large, complex application). What is it? Be specific, and describe it in a paragraph or two, with an illustration.

Q5 (0.5*4=2 points). A highly effective, powerful and conceptually simple way to analyze data, is to employ 'dataflow'. Name two uses of dataflow that you learned from the course (we covered three).

a.

b.

Data can also be processed in multiple stages (pipelining). Name two architectures we covered, that permit such multistage data analysis.

a.

b.


Q6 (3 points). What is the most flexible way to model ('any') data? And, what structure can be used to do so?

Q7 (2 points). Functional programming is an expressive, compact way of specifying data processing operations. You have seen two examples in the course - what are they?

a.

b.

Q8 (2+2=4 points). Represent the following graph in (valid!) JSON notation [suitable for storing in a text file, and reading it back to construct the graph] - do it two different ways!



Q9 (0.5*10=5 points). An autonomous car is deployed on the streets, having its neural networks trained (using tens of thousands of hours of hand-labeled traffic videos). 'Autonomous warfare' is a scenario that is increased being discussed by analysts, where in a (distant?!) future, self-guided soldier 'bots' would conduct war operations against 'live' enemies (who presumably don't have the means to deploy automated soldiers). Our auto-guided soldier bots could operate on land (with wheels or limbs), in the water, and be airborne as well ('killer drones').

What would the required training data be? Think broadly (including multi-sensory modes!), and name 10 different classes (detection targets, ie. things to learn to recognize) that would be useful. In other words, what can we teach our bots' neural nets?

- 1.**
- 2.**
- 3.**
- 4.**
- 5.**
- 6.**
- 7.**
- 8.**
- 9.**
- 10.**

Q10 (5 points). Many events/phenomena in our lives have a spatial dependence - eg. if you live under the flight path of airplanes (eg. in Westchester, near LAX), your car would need to get cleaned quite often because of black, oily jet fuel getting deposited on it by landing planes! To prove such a causal link (landing airplanes cause oily buildups), you'd need to plot on a map: flight paths, and amount of soot (for example) per week found on cars near and away from the flight path. Cars closer to, and on the flight path, would show higher levels. Given the following list of causes and effects in no particular order, pick out 5 pairs of (cause, effect) relations that can be investigated for possible occurrence (or merely shown/documented), by plotting data on a map. Eg. the above example would be listed as (flight path, oily deposit). Note - your pairings need to be plausible, not frivolous.

Here is the list: flight path, freeways, particulate matter, power lines, noise, cancer, known gang locations, oily deposit, cellphone towers, wealthy neighborhoods, graffiti, BMW dealerships, low-income neighborhoods, crime, fast-food restaurants, vibrations (rattling), banks, pawn shops.

Among other things, such visualizations can be used to effect public policy, document social disparities, etc.

- a.**
- b.**
- c.**
- d.**
- e.**

Bonus (1 point). What is being indicated below?

